

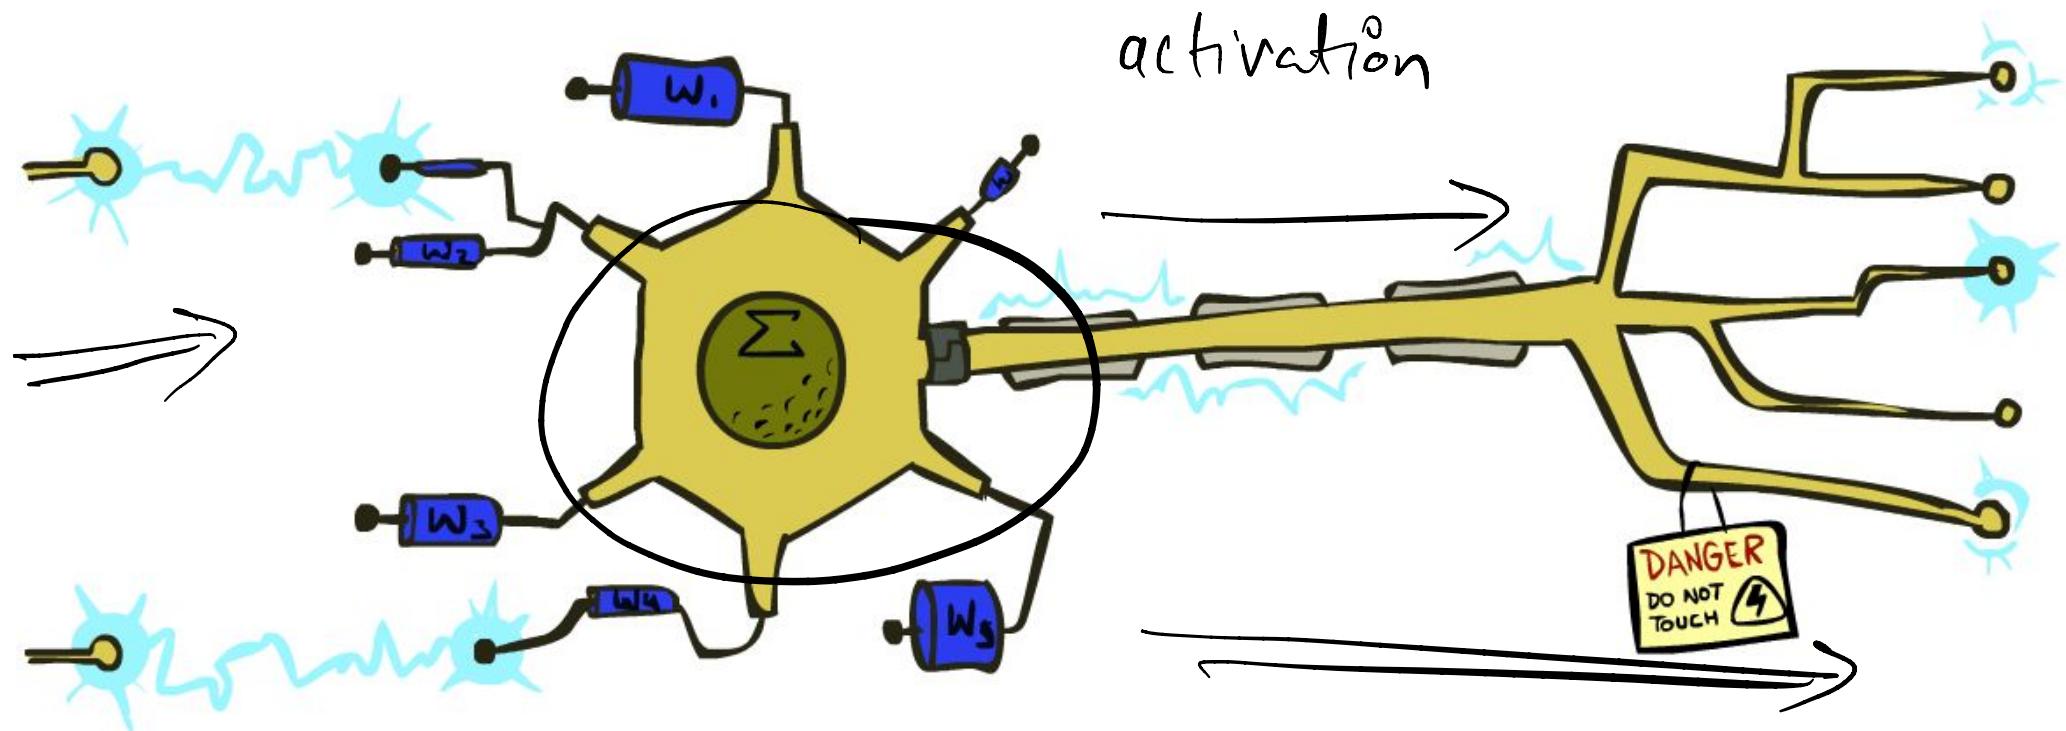
Perceptrons and Logistic Regression

neural network

deep learning

Machine learning

Perceptrons and Logistic Regression

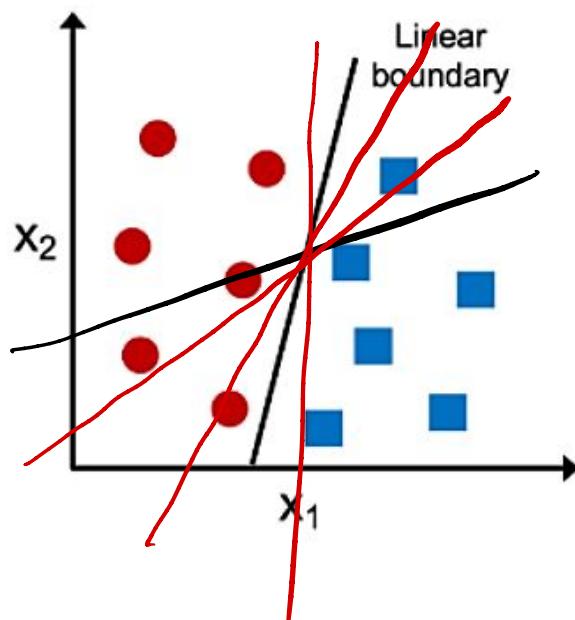


Classifiers

equation \Rightarrow parameters
if
weights and bias

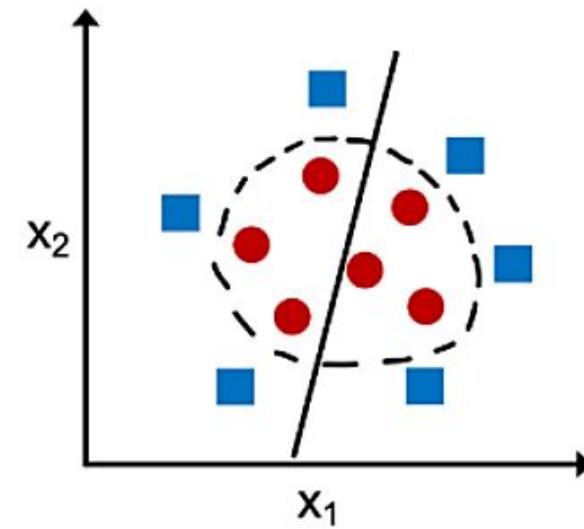
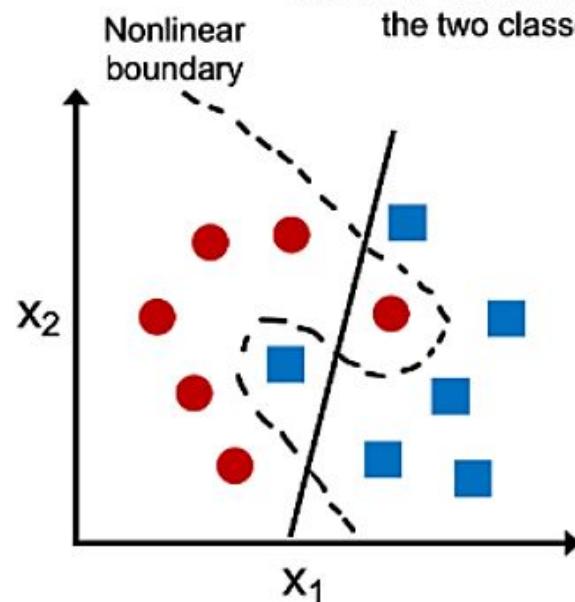
Linearly separable

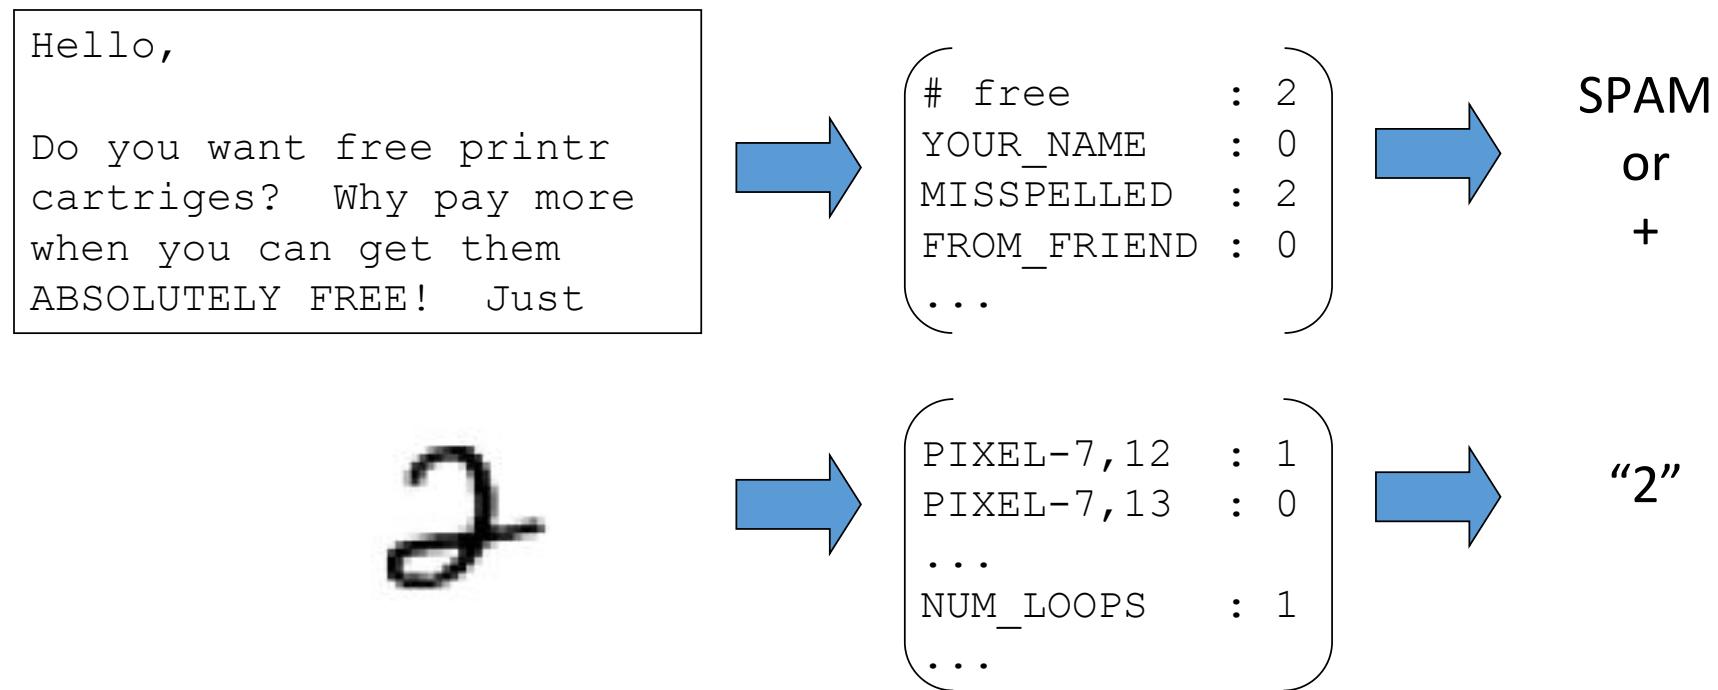
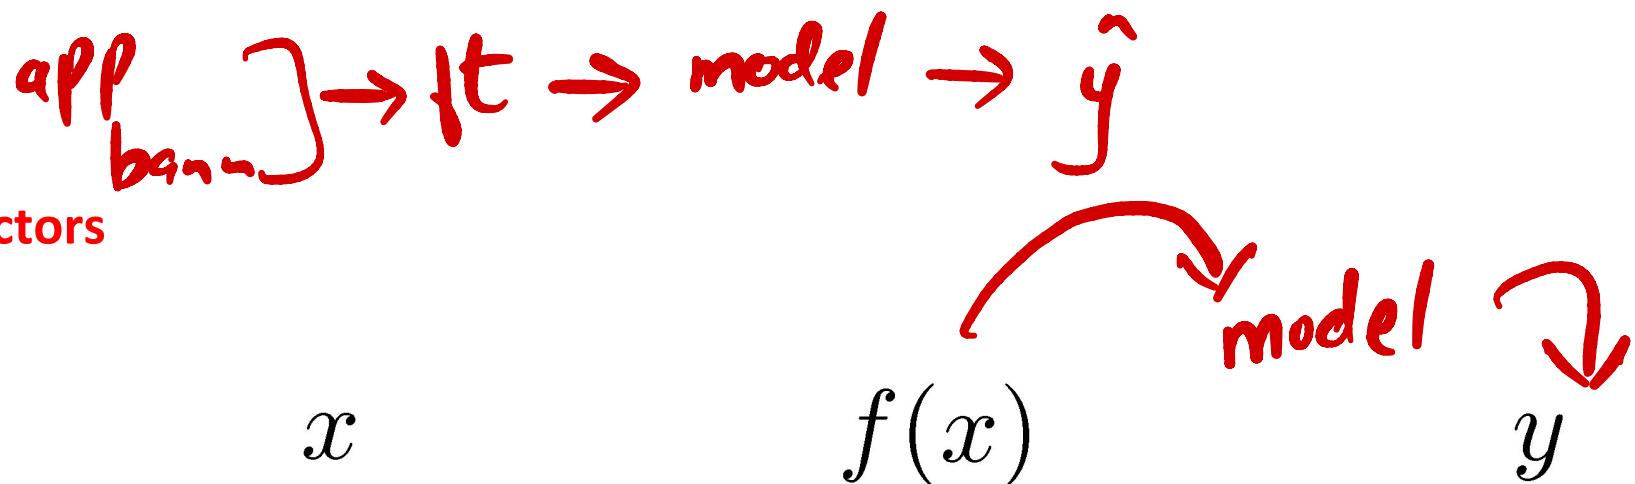
A linear decision boundary that separates the two classes exists



Not linearly separable

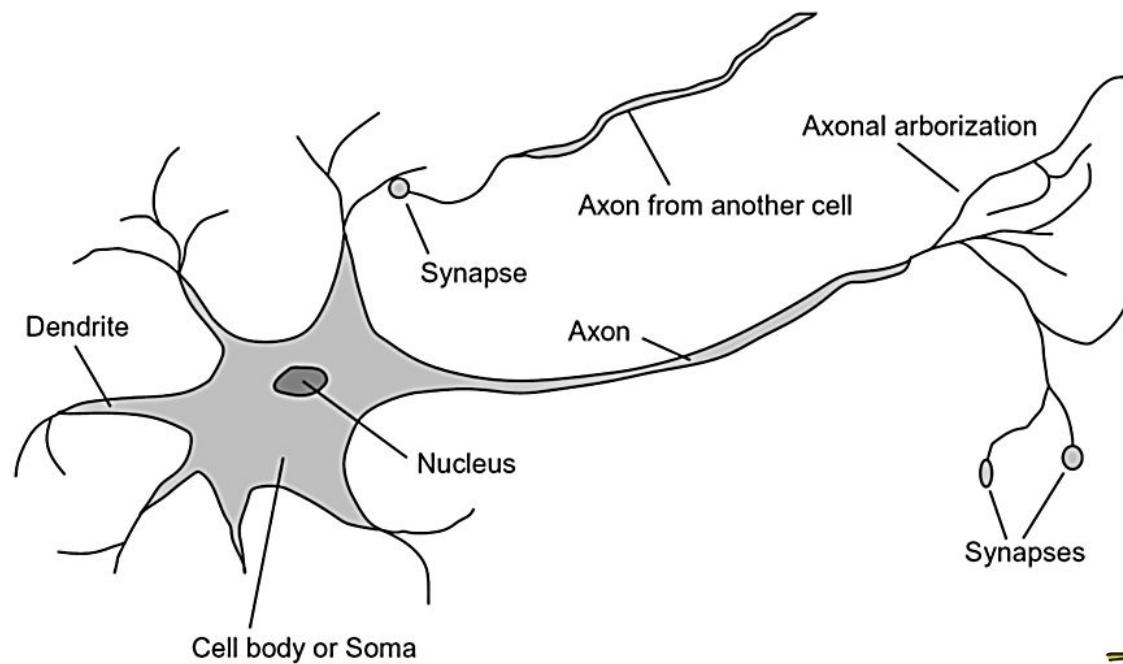
No linear decision boundary that separates the two classes perfectly exists





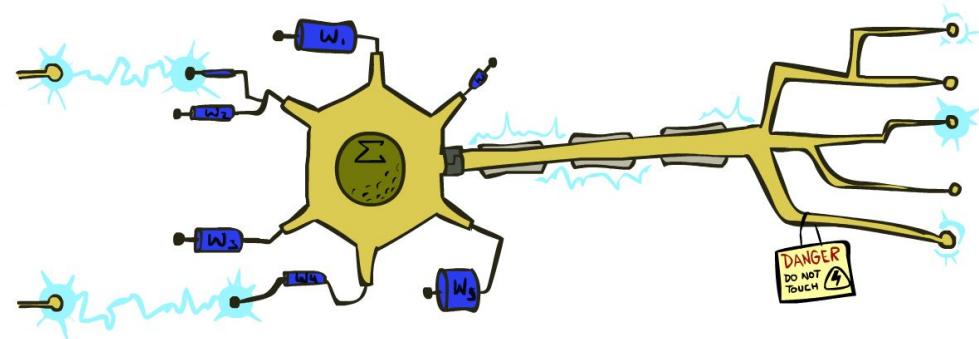
Some (Simplified) Biology

- Very loose inspiration: human neurons



IP → O → O → O → O → O → class' fitted ~

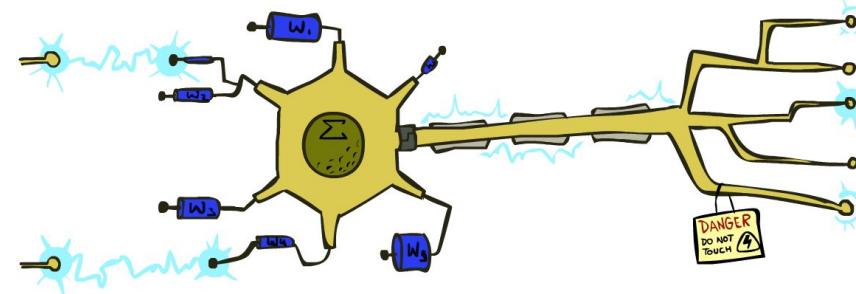
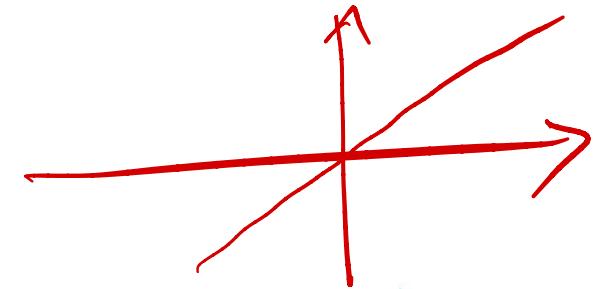
Perception is
a single neuron



Linear Classifiers

- Inputs are **feature values**
- Each feature has a **weight**
- Sum is the **activation**

Sigmoid , relu

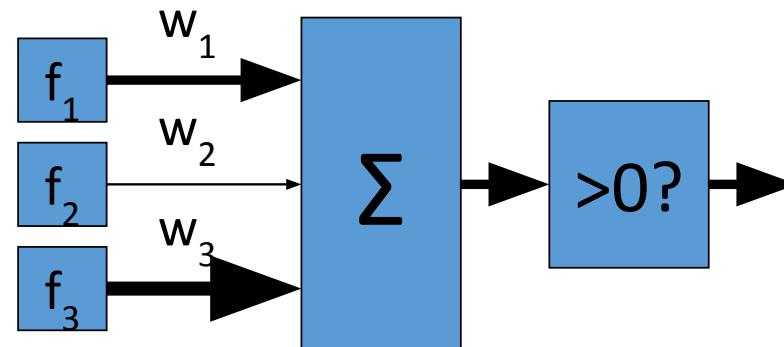


$$\text{activation}_w(x) = \sum_i w_i \cdot f_i(x) = w \cdot f(x)$$

↳ feature from a

- If the activation is:

- Positive, output +1
- Negative, output -1



all example

Weights

- Binary case: compare features to a weight vector
- Learning: figure out the weight vector from examples

$\sum f(x) w \rightarrow$ we
-ve
 update weights

# free	:	4
YOUR_NAME	:	-1
MISSPELLED	:	1
FROM_FRIEND	:	-3
...		

w

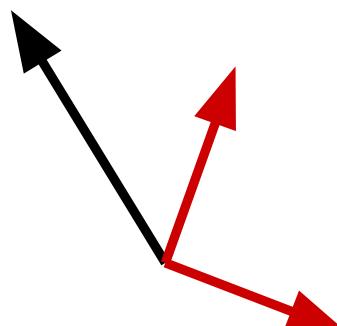
$f(x_1)$

# free	:	2
YOUR_NAME	:	0
MISSPELLED	:	2
FROM_FRIEND	:	0
...		

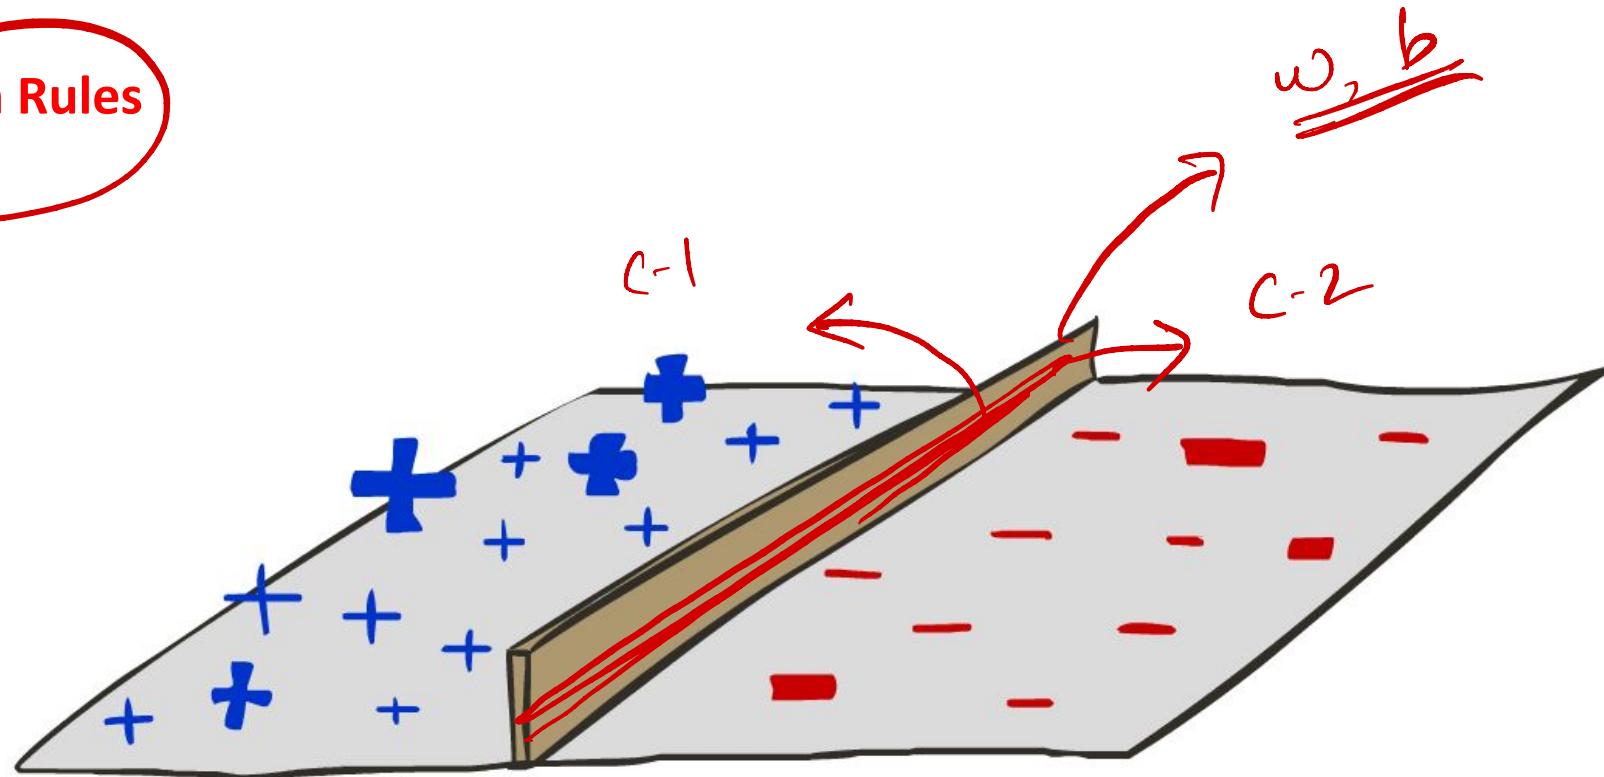
$f(x_2)$

# free	:	0
YOUR_NAME	:	1
MISSPELLED	:	1
FROM_FRIEND	:	1
...		

Dot product $w \cdot f$ positive
 means the positive class



Decision Rules

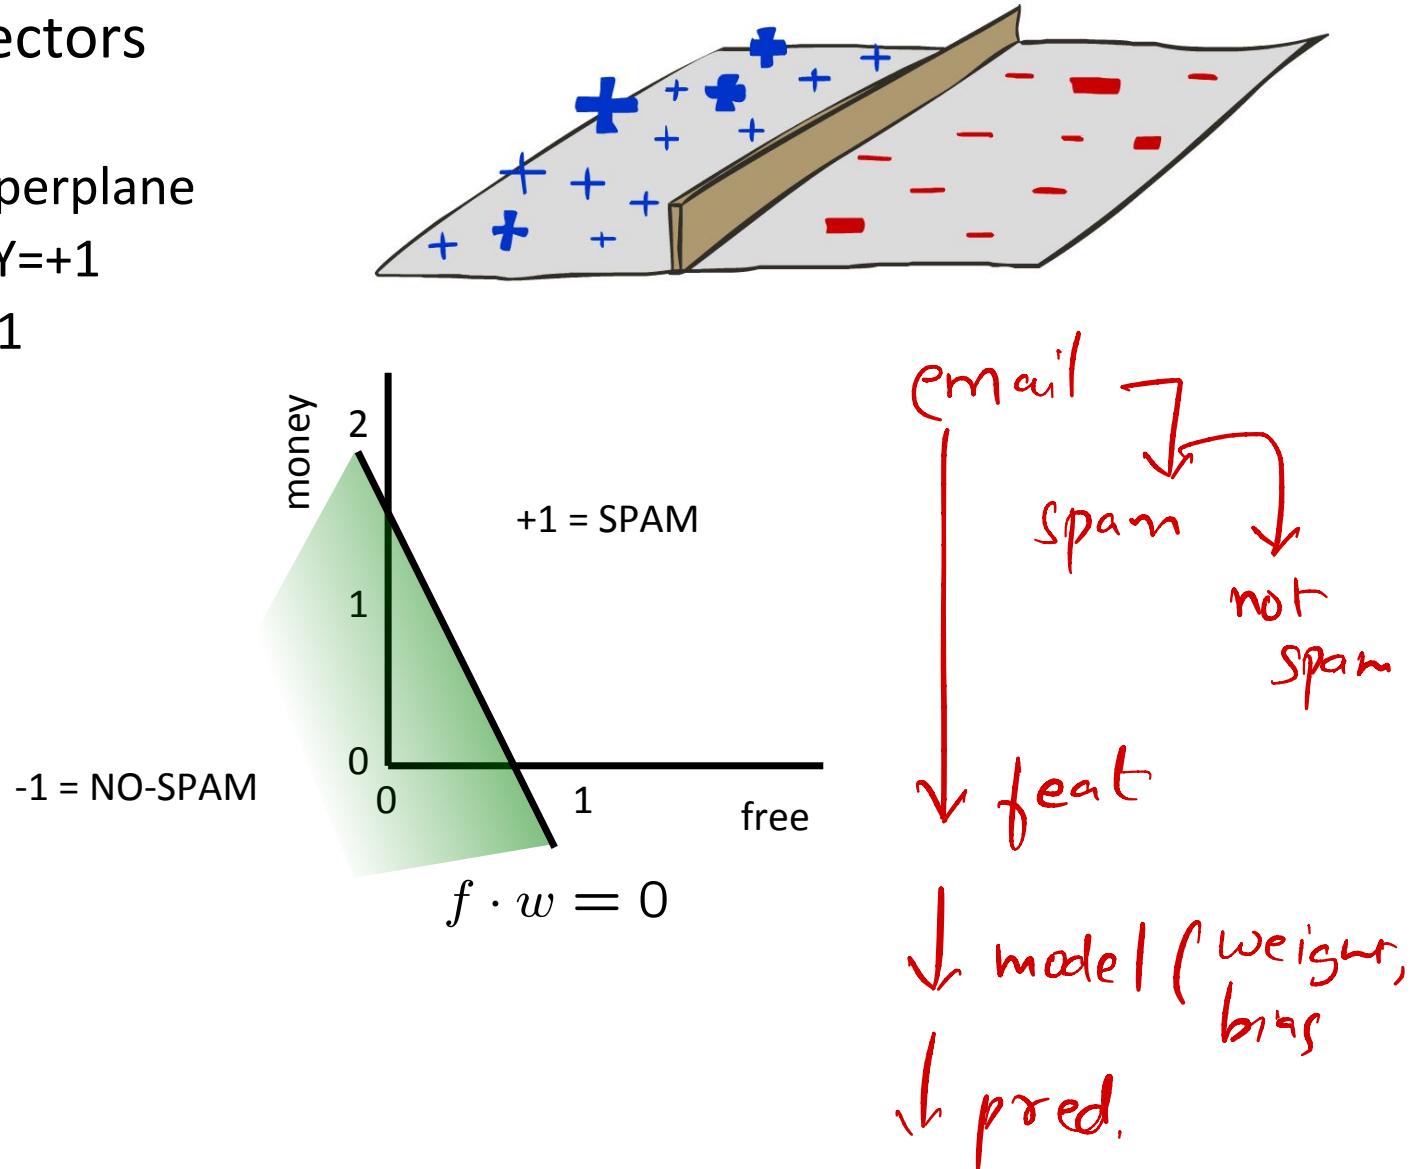


true value
predicted value

Binary Decision Rule

- In the space of feature vectors
 - Examples are points
 - Any** weight vector is a hyperplane
 - One side corresponds to $Y=+1$
 - Other corresponds to $Y=-1$

w
BIAS : -3
free : 4
money : 2
...

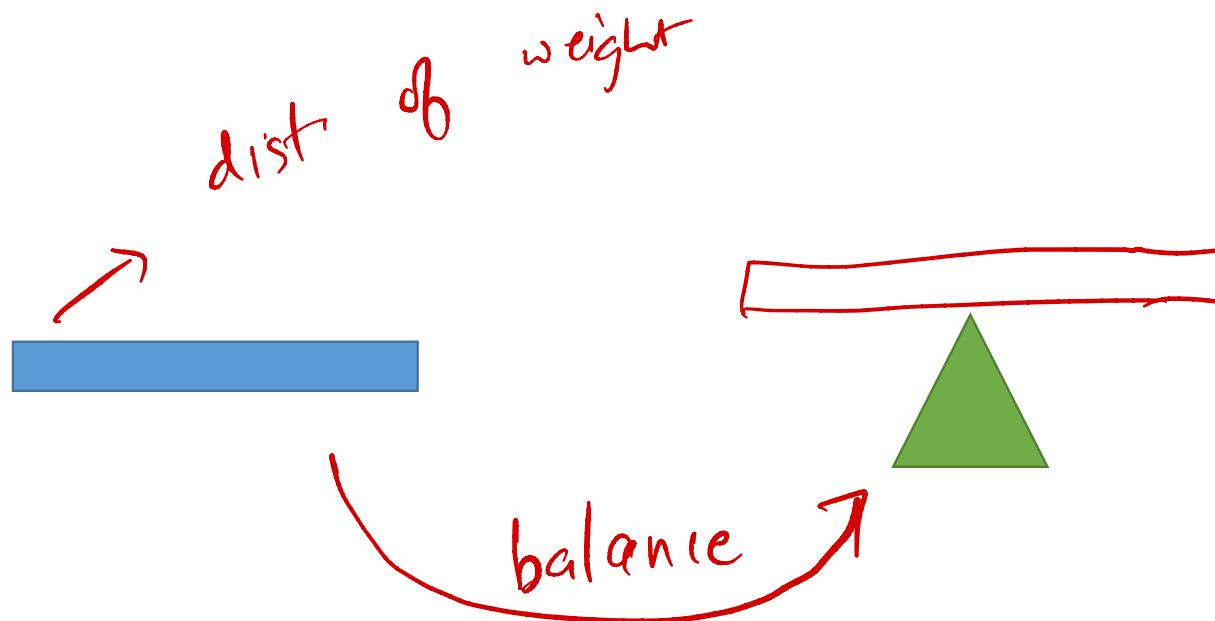


$$\omega = \omega + f \cdot y^*$$

Weight Updates

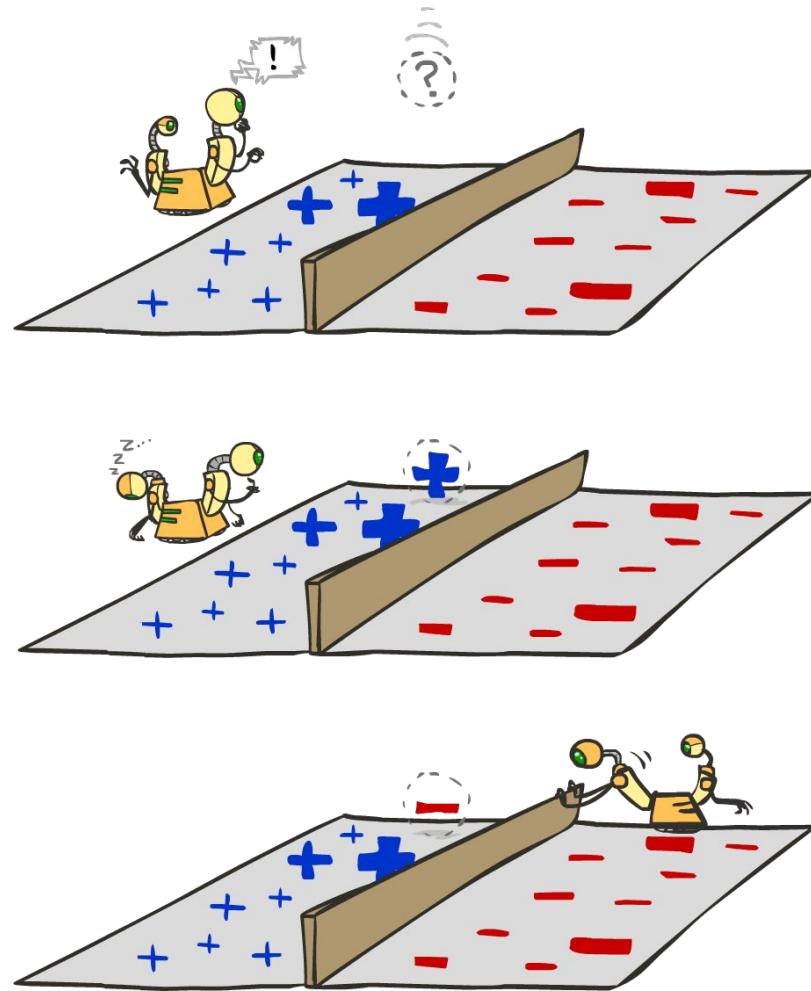
Example: to balance

Shishir Maheshwari



Learning: Binary Perceptron

- Start with weights = 0
- For each training instance:
 - Classify with current weights
 - If correct (i.e., $y=y^*$), no change!
 - If wrong: adjust the weight vector



McCulloch Pitts model (MP model)

→ 1949

↳ AND, OR, XOR, NOT, ...

$$w \in [-1, 1]$$

$$\begin{aligned} z &= \sum x_i w_i \\ \hat{y} &= \begin{cases} 1, & z \geq 0 \\ 0, & z < 0 \end{cases} \end{aligned}$$

↙ true value

x_1	x_2	y
0	0	0
0	1	0
1	0	0
1	1	1

$$z = \sum_{i=1}^n x_i^\circ w_i^\circ, \quad \hat{y} = \begin{cases} 1, & z \geq \theta \\ 0, & z < \theta \end{cases}$$

AND

x_1	x_2	y
0	0	0
0	1	0
1	0	0
1	1	1

find w, θ

$$\hat{y} = y$$

$$w \in [-1, 1]$$

\Rightarrow result, w_1 and w_2

Initialise w

$\rightarrow z, \hat{y}$

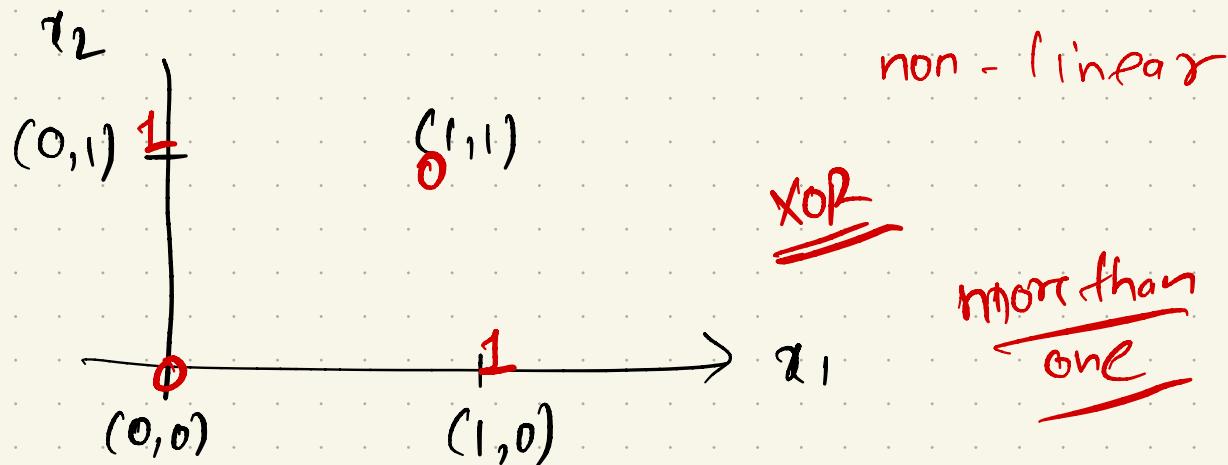
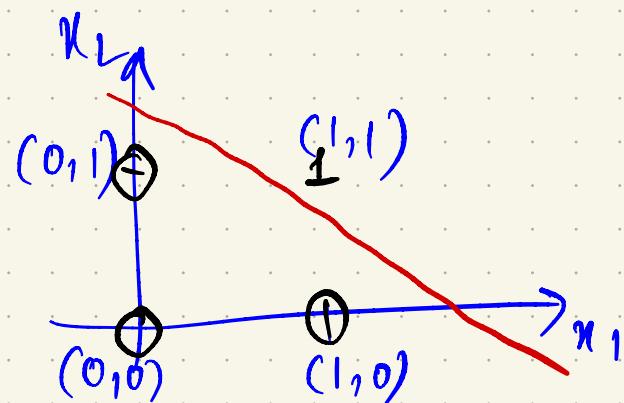
\rightarrow find θ such that $y = \hat{y}$

$$w_1 = 1 \quad w_2 = 1, \quad \begin{matrix} 0 \rightarrow 0 \\ 1 \rightarrow 0 \\ 1 \rightarrow 0 \\ 2 \rightarrow 1 \end{matrix} \quad \theta = 2 \quad \hat{y} = y$$

$$w_1 = -1, \quad w_2 = 1, \quad \begin{matrix} 0 \\ -1 \\ 1 \\ 0 \end{matrix} \quad \theta = ? \quad \hat{y} = y$$

XOR

x_1	x_2	y	$w_1 = ?$	$w_2 = ?$	$\Rightarrow \theta = ? \Rightarrow$	$\hat{y} = y$
0	0	0				
0	1	1				
1	0	1				
1	1	0				



issue with MP model ?

\Rightarrow no auto update

$\Rightarrow w_1, b$, or any other \Rightarrow auto update

Learning: Binary Perceptron

- Start with weights = 0
- For each training instance:
 - Classify with current weights

$$\hat{y} = \begin{cases} +1 & \text{if } w \cdot f(x) \geq 0 \\ -1 & \text{if } w \cdot f(x) < 0 \end{cases}$$

- If correct (i.e., $\hat{y}=y^*$), no change!
- If wrong: adjust the weight vector by adding or subtracting the feature vector.
Subtract if y^* is -1.

$$w = w + y^* \cdot f \Rightarrow \underline{\text{update rule}}$$

\hookrightarrow true label

$\hat{y} \neq y^* \Rightarrow \text{update}$
 $\hat{y} = y^* \Rightarrow \text{no update}$

①	②	③	④	⑤	⑥
$x = -1$	-2	-3	1	2	3
$y^* = -1$	-1	-1	1	1	1

$w = \text{initialise}$, $y = ?$

$x == f(x)$

achieve?

(1) $w = -3$

(2) $y =$ for each

(3) update

$y_1 = 3 \Rightarrow 1 \neq -1$

$$\begin{aligned} w &= (-3) + (-1)(-1) \\ &= -2 \end{aligned}$$

$y_2 = 4 \Rightarrow 1 \neq -1$

$$w = (-2) + (-2)(-1) = 0$$

$y_3 = 0 \Rightarrow 1 \neq -1$

$$w = (0) + (3) = 3$$

$y_4 = \checkmark$

$y_5 = \checkmark$

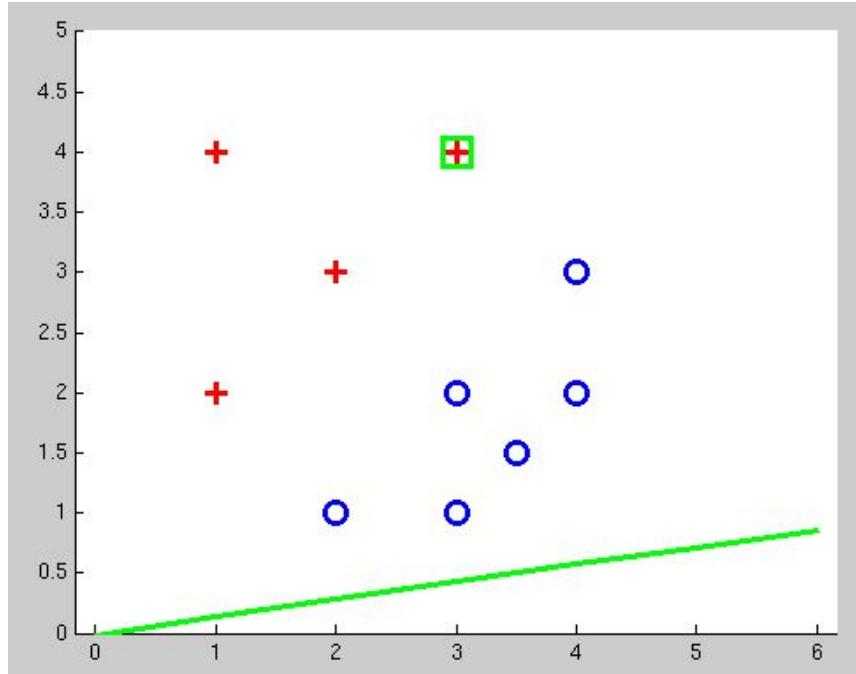
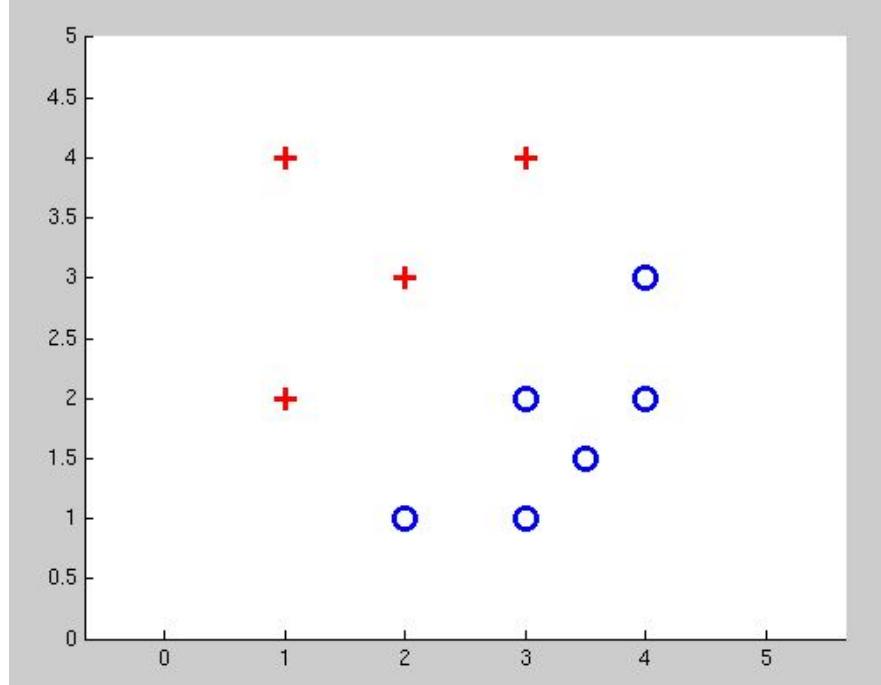
$y_6 = \checkmark$

$y_1 = \checkmark$

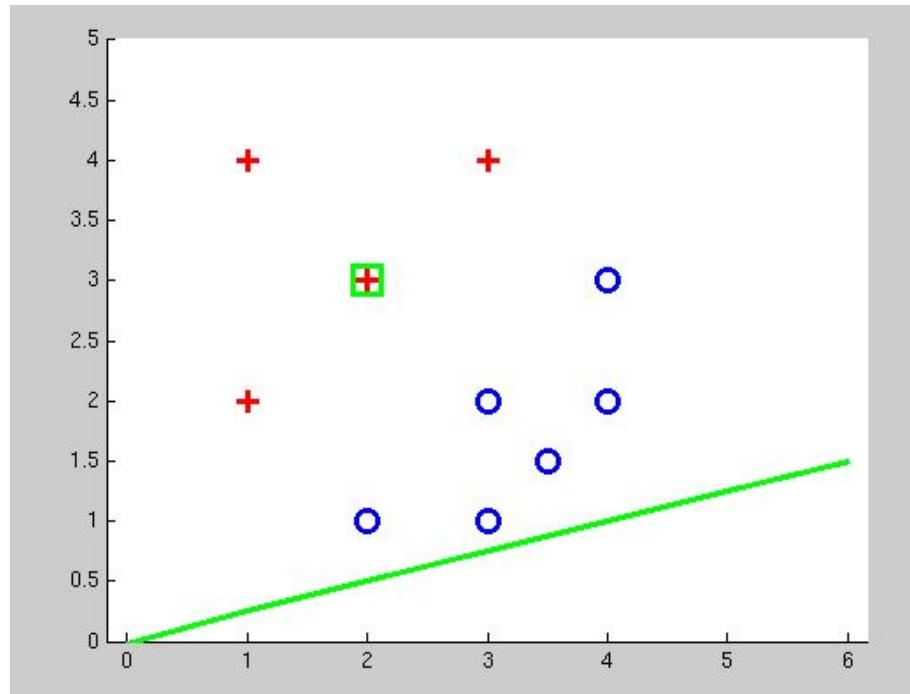
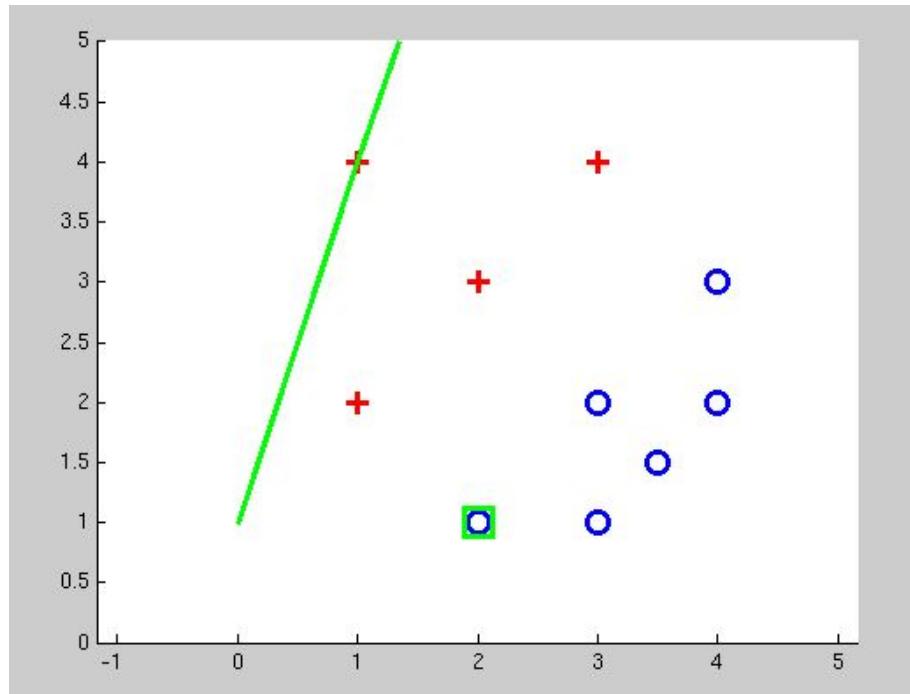
$y_2 = \checkmark$

$y_3 = \checkmark$

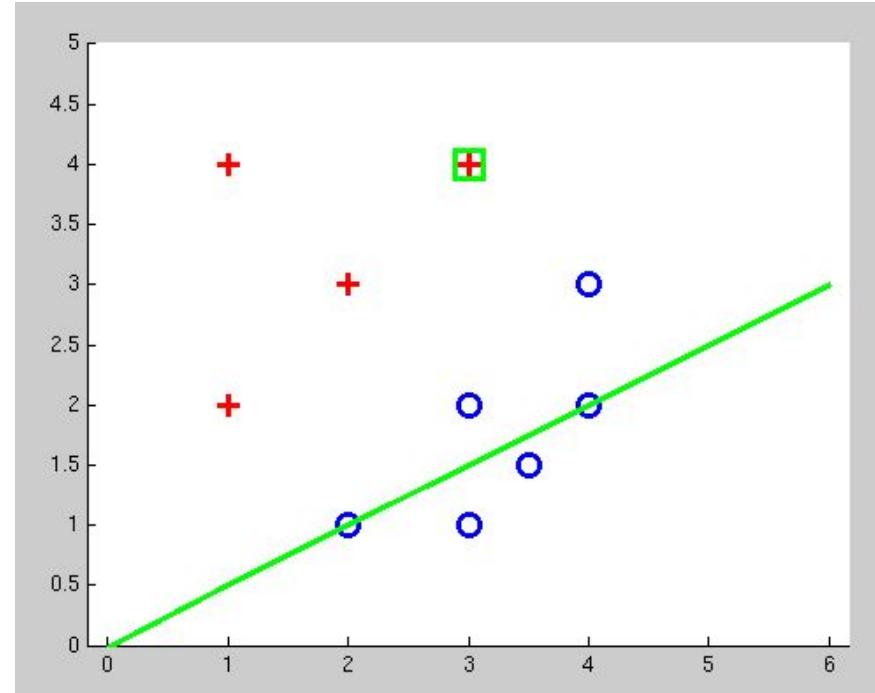
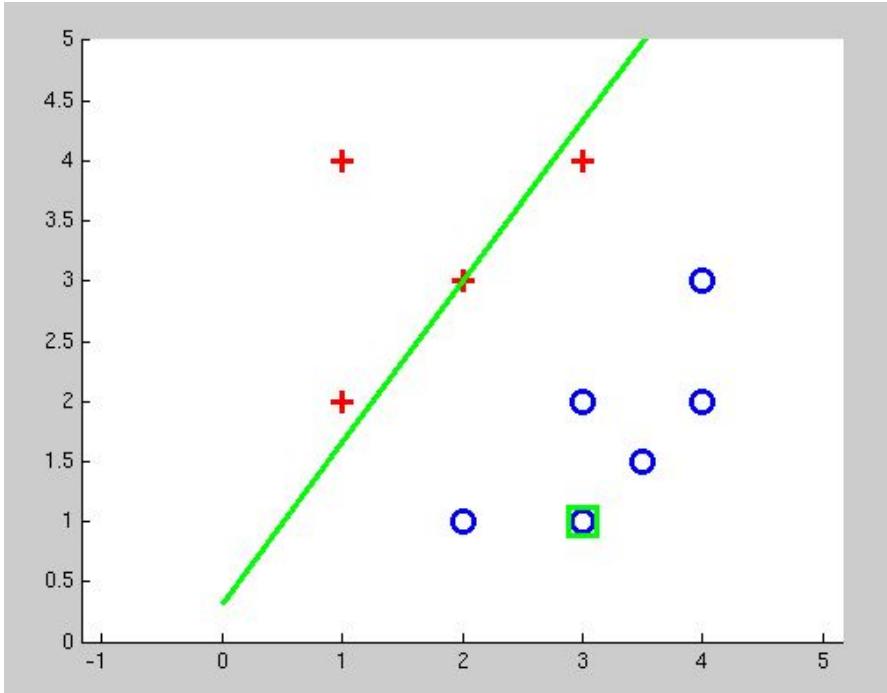
Examples: Perceptron



Examples: Perceptron

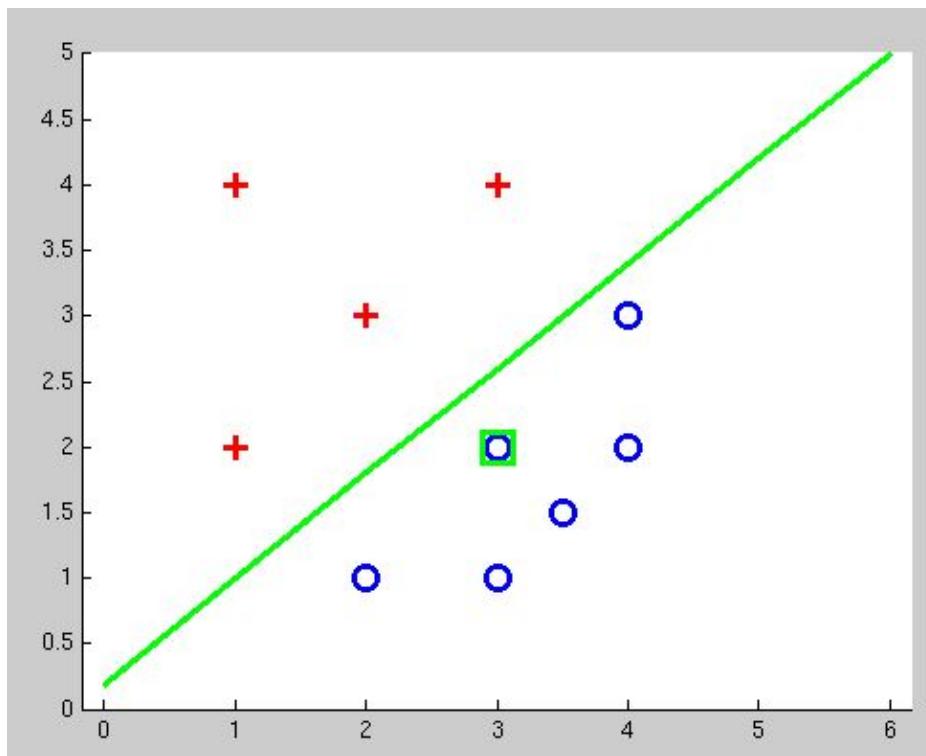


Examples: Perceptron



Examples: Perceptron

all input \Rightarrow one epoch
many epochs \Rightarrow find ans



$$x = \begin{matrix} 10 & 11 & 12 & 1 & 2 & 3 \end{matrix}$$

$$y^* = \begin{matrix} -1 & -1 & -1 & 1 & 1 & 1 \end{matrix}$$

$w = -3$, find $w = ?$ for $y = y^*$

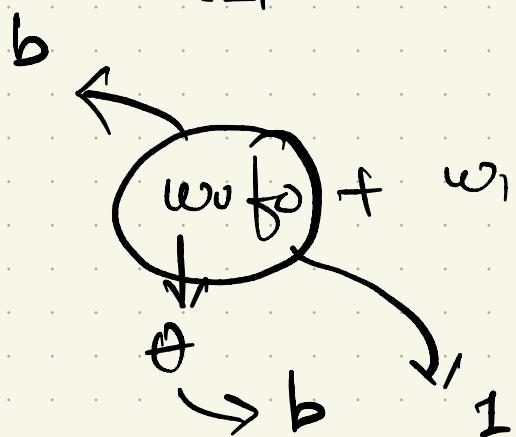
$$w = -1/-2/-3$$

$$b = 5/6/8$$

bias = ?

$$\sum_{i=1}^n w_i f_i(x) > \theta \Rightarrow$$

$$\sum_{i=1}^n w_i f_i + \theta = 0 \Rightarrow$$



$$= \boxed{\sum_{i=0}^n w_i f_i = y_{pred}}$$

$$y_{\text{pred}} = b + \sum_{i=1}^N w_i f_i(x)$$

$$\sum_{i=0}^N w_i f_i(x)$$

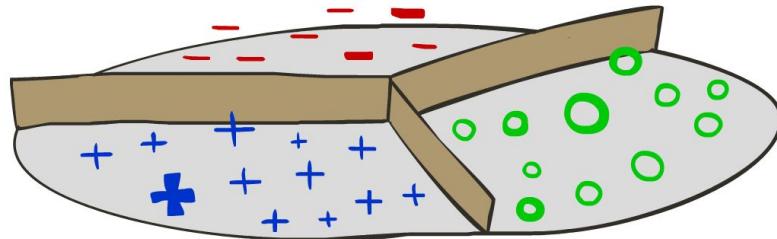
$$\boxed{\begin{aligned} w &= w + y^* t \\ b &= b + y^* \end{aligned}}$$

2 - binary

Multiclass Decision Rule

- If we have multiple classes:
 - A weight vector for each class:

$$w_y$$



- Score (activation) of a class y :

$$w_y \cdot f(x)$$

- Prediction highest score wins

$$y = \arg \max_y [w_y \cdot f(x)]$$

3 - class c_1 c_2 c_3
 w_{c_1} w_{c_2} w_{c_3}

→ initial value of w

→ $f(a), w$

→ $y = \arg \max [w, f(a)]$

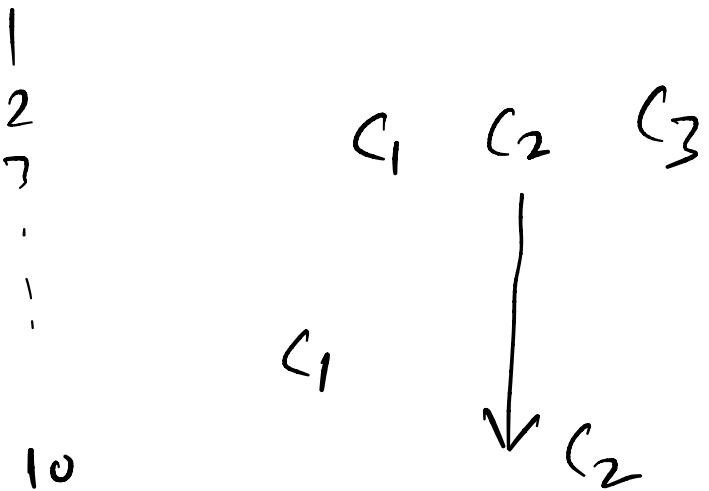
Learning: Multiclass Perceptron

- Start with all weights = 0
- Pick up training examples one by one
- Predict with current weights

$$y = \arg \max_y w_y \cdot f(x)$$

- If correct, no change!
- If wrong: lower score of wrong answer, raise score of right answer

$$\left[\begin{array}{l} w_y = w_y - f(x) \\ w_{y^*} = w_{y^*} + f(x) \end{array} \right]$$



$c_1 \quad c_2 \quad c_3$
 $\geq 0 \quad c_1$
 $< 0 \quad c_2$

true = = pred \Rightarrow no update
 $f \Rightarrow$ update

$$x_1 \rightarrow f(x_1) \Rightarrow f(x_1) = \begin{bmatrix} 1 \\ \text{win}=1 \\ \text{game}=0 \\ \text{vote}=1 \\ \text{the}=1 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 0 \\ 1 \\ 1 \end{bmatrix}$$

data \Rightarrow feature
 $x_1 \Rightarrow f(x_1)$

Example: Multiclass Perceptron

(1) "win the vote" $\Rightarrow f(x) \Rightarrow$ $f(x) \Rightarrow \max \Rightarrow =$ no update

(2) "win the election"

(3) "win the game"

w_{SPORTS} initial value

BIAS : 1
win : 0
game : 0
vote : 0
the : 0
... : ...

$w_{POLITICS}$

BIAS : 0
win : 0
game : 0
vote : 0
the : 0
... : ...

w_{TECH}

BIAS : 0
win : 0
game : 0
vote : 0
the : 0
... : ...

$$f(x_1) = \begin{bmatrix} 1 \\ -1 \\ 0 \\ 1 \\ 1 \end{bmatrix}$$

$$\omega_{\text{sp}} = \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

$$\omega_{\text{poli}} = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

$$\omega_{\text{tech}} = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

$x_1 \rightarrow$ sports
predicted
true label \Rightarrow politics

$$\Rightarrow \begin{array}{c} f(x_1) \cdot \omega_{\text{sp}} \\ 1 \times \\ \downarrow \end{array} \quad \begin{array}{c} f(x_1) \cdot \omega_{\text{poli}} \\ 0 \times \\ \uparrow \end{array} \quad \begin{array}{c} f(x_1) \cdot \omega_{\text{tech}} \\ 0 \\ nc \end{array}$$

$$\begin{aligned} \omega_{\text{sp}} &= \omega_{\text{sp}} - \\ &= \begin{bmatrix} 0 \\ -1 \\ 0 \\ -1 \\ -1 \end{bmatrix} \end{aligned}$$

$$\begin{aligned} \omega_{\text{poli}} &= \omega_{\text{poli}} + \\ \omega_{\text{poli}} &= \begin{bmatrix} 1 \\ -1 \\ 0 \\ 1 \\ 1 \end{bmatrix} \end{aligned}$$

$$f(x_2) = \begin{bmatrix} 1 \\ -2 \\ 0 \\ 0 \end{bmatrix} \quad w_{sp} \quad w_{poli} \quad w_{tech}$$

$f(x_2) \rightarrow \text{politics}$
true label \rightarrow poli

$$f(x_3) = \begin{bmatrix} 1 \\ -2 \\ 1 \\ 0 \end{bmatrix} \quad w_{sp} \quad w_{poli} \quad w_{tech}$$

$f(x_3) \rightarrow \text{polit}$
true \rightarrow sports

$$w_{sp} = w_{sp} +$$

$$w_{poli} = w_{poli} -$$

$$w_{sp} = \begin{bmatrix} 1 \\ 0 \\ -1 \\ 0 \end{bmatrix}$$

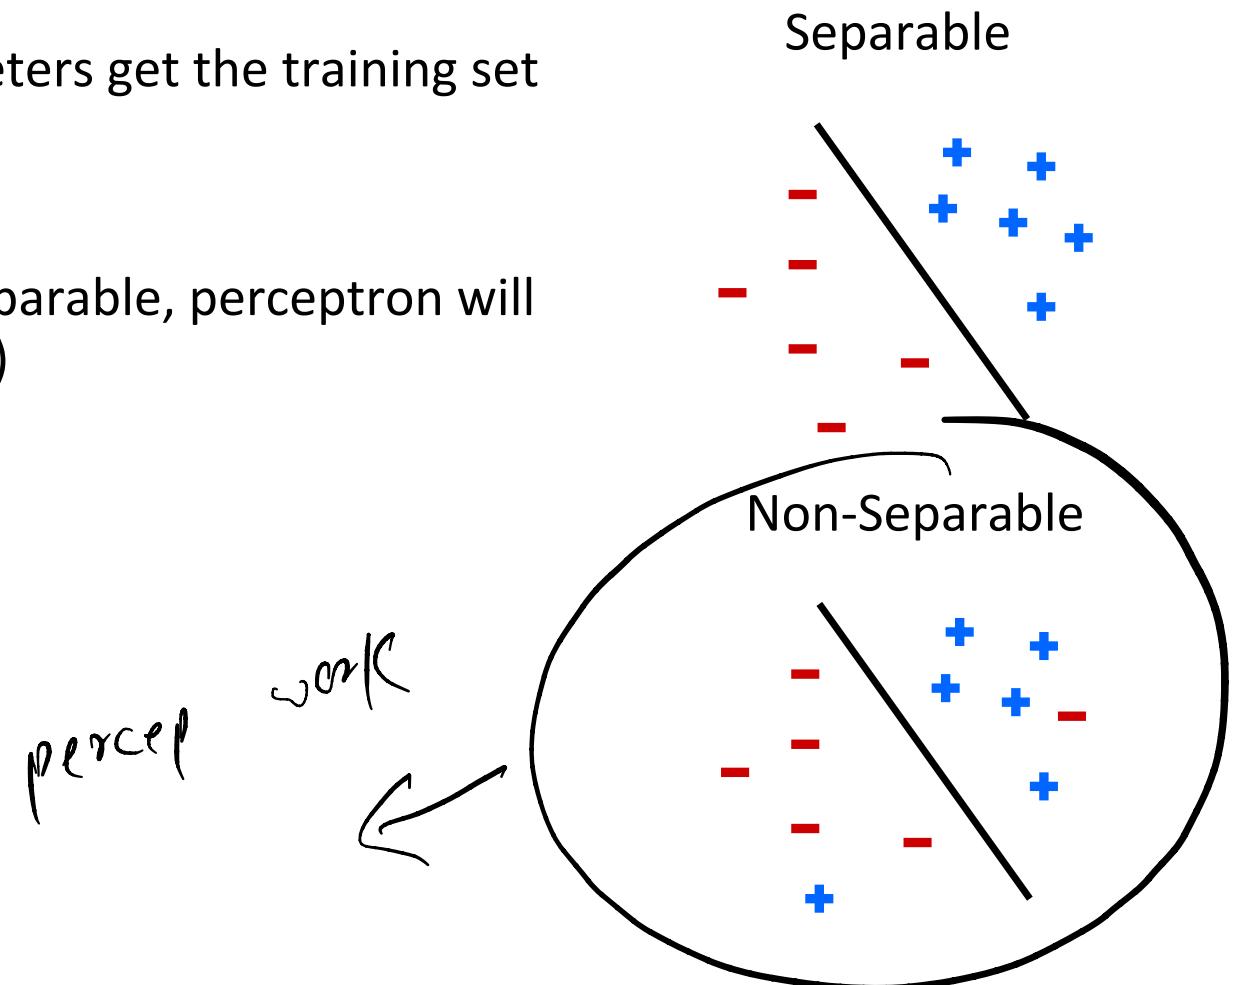
$$w_{poli} = \begin{bmatrix} 0 \\ -1 \\ 1 \\ 0 \end{bmatrix}$$

$$w_{tech} = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

- init ω (given / random)
 - find score
 - from score find pred. class
- \Rightarrow pred == true , no update
- \Rightarrow \neq update

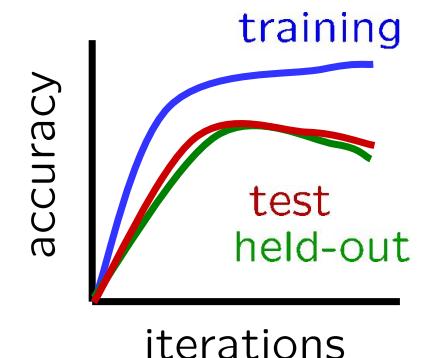
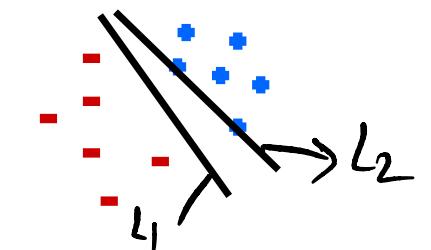
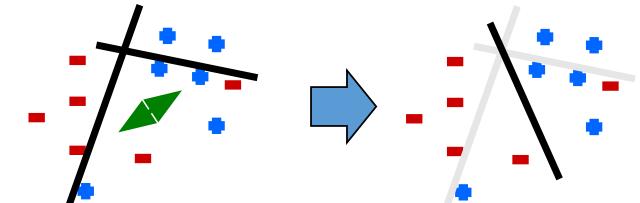
Properties of Perceptrons

- Separability: true if some parameters get the training set perfectly correct
- Convergence: if the training is separable, perceptron will eventually converge (binary case)



Problems with the Perceptron

- Noise: if the data isn't separable, weights might thrash
 - Averaging weight vectors over time can help (averaged perceptron)
- Mediocre generalization: finds a “barely” separating solution
- Overtraining: test / held-out accuracy usually rises, then falls
 - Overtraining is a kind of overfitting



Improving the Perceptron

mapp i/p → o/p []